

```
# =====
# 1. Data Loading
# =====

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix

# Load dataset
df = pd.read_csv("/content/drive/MyDrive/DS160/Final Project/AI_Impact_on_Jobs_2030.csv")

df.head()
```

	Job_Title	Average_Salary	Years_Experience	Education_Level	AI_Exposure_Index	Tech_Growth_Factor	Au
0	Security Guard	45795	28	Master's	0.18	1.28	
1	Research Scientist	133355	20	PhD	0.62	1.11	
2	Construction Worker	146216	2	High School	0.86	1.18	
3	Software Engineer	136530	13	PhD	0.39	0.68	
4	Financial Analyst	70397	22	High School	0.52	1.46	

```
# Check dataset info
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Job_Title        3000 non-null   object  
 1   Average_Salary   3000 non-null   int64   
 2   Years_Experience 3000 non-null   int64   
 3   Education_Level  3000 non-null   object  
 4   AI_Exposure_Index 3000 non-null   float64 
 5   Tech_Growth_Factor 3000 non-null   float64 
 6   Automation_Probability_2030 3000 non-null   float64 
 7   Risk_Category    3000 non-null   object  
 8   Skill_1          3000 non-null   float64 
 9   Skill_2          3000 non-null   float64 
 10  Skill_3          3000 non-null   float64 
 11  Skill_4          3000 non-null   float64 
 12  Skill_5          3000 non-null   float64 
 13  Skill_6          3000 non-null   float64 
 14  Skill_7          3000 non-null   float64 
 15  Skill_8          3000 non-null   float64 
 16  Skill_9          3000 non-null   float64 
 17  Skill_10         3000 non-null   float64 
dtypes: float64(13), int64(2), object(3)
memory usage: 422.0+ KB
```

```
# Check shape
```

```
df.shape
```

```
(3000, 18)
```

```
df.isnull().sum()
```

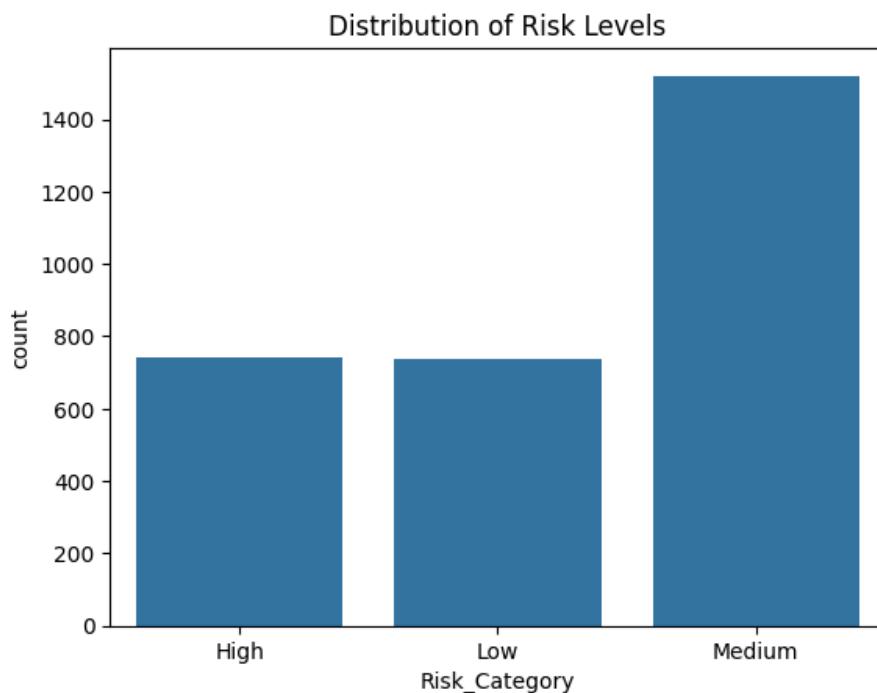
	0
Job_Title	0
Average_Salary	0
Years_Experience	0
Education_Level	0
AI_Exposure_Index	0
Tech_Growth_Factor	0
Automation_Probability_2030	0
Risk_Category	0
Skill_1	0
Skill_2	0
Skill_3	0
Skill_4	0
Skill_5	0
Skill_6	0
Skill_7	0
Skill_8	0
Skill_9	0
Skill_10	0

```
dtype: int64
```

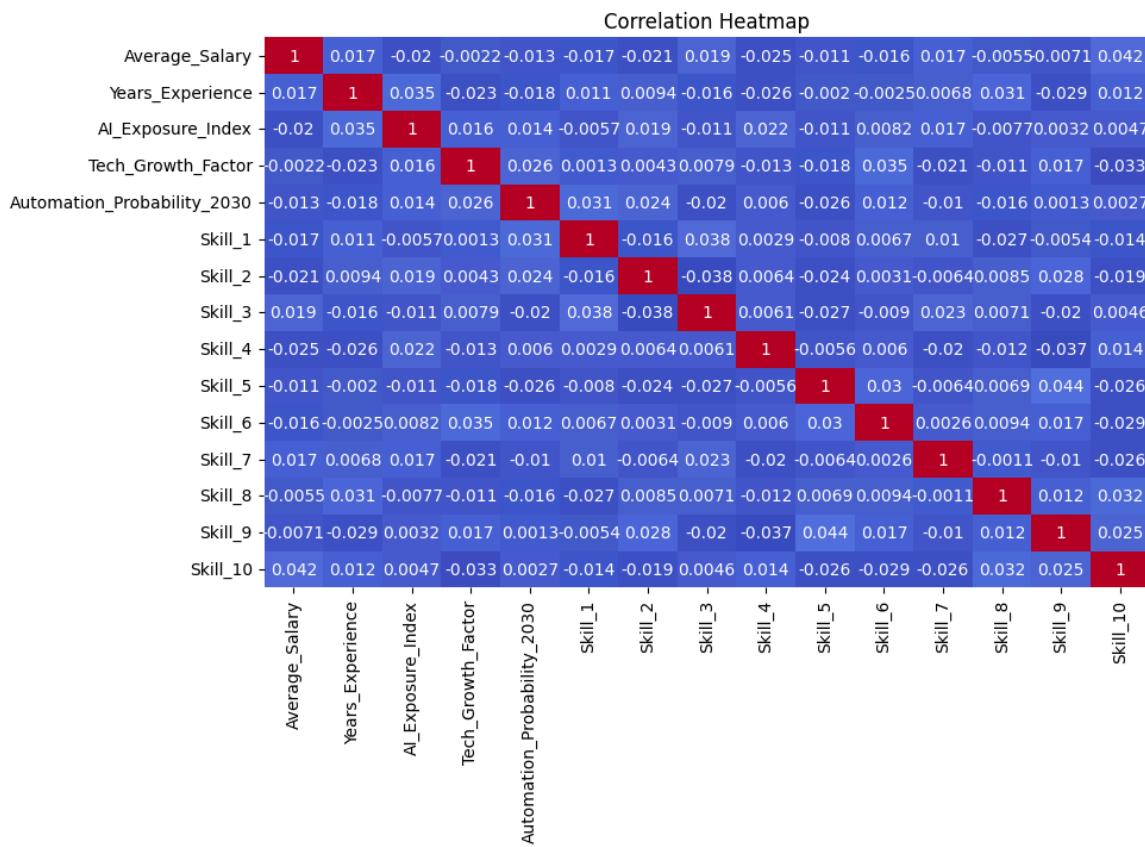
```
df.describe(include="all")
```

	Job_Title	Average_Salary	Years_Experience	Education_Level	AI_Exposure_Index	Tech_Growth_Factor
count	3000	3000.000000	3000.000000	3000	3000.000000	3000.000000
unique	20	NaN	NaN	4	NaN	NaN
top	Software Engineer	NaN	NaN	High School	NaN	NaN
freq	175	NaN	NaN	784	NaN	NaN
mean	NaN	89372.279000	14.677667	NaN	0.501283	0.995343
std	NaN	34608.088767	8.739788	NaN	0.284004	0.287669
min	NaN	30030.000000	0.000000	NaN	0.000000	0.500000
25%	NaN	58640.000000	7.000000	NaN	0.260000	0.740000
50%	NaN	89318.000000	15.000000	NaN	0.500000	1.000000
75%	NaN	119086.500000	22.000000	NaN	0.740000	1.240000
max	NaN	149798.000000	29.000000	NaN	1.000000	1.500000

```
sns.countplot(data=df, x="Risk_Category")
plt.title("Distribution of Risk Levels")
plt.show()
```



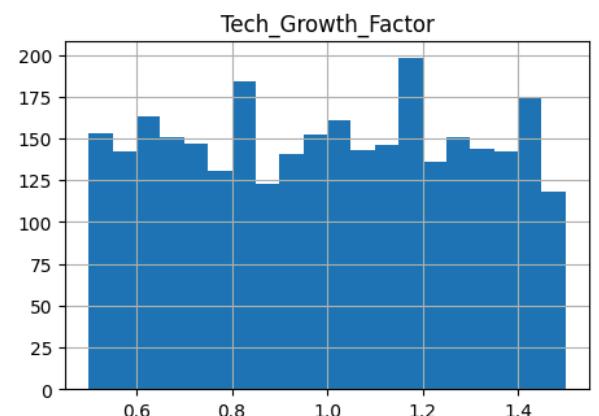
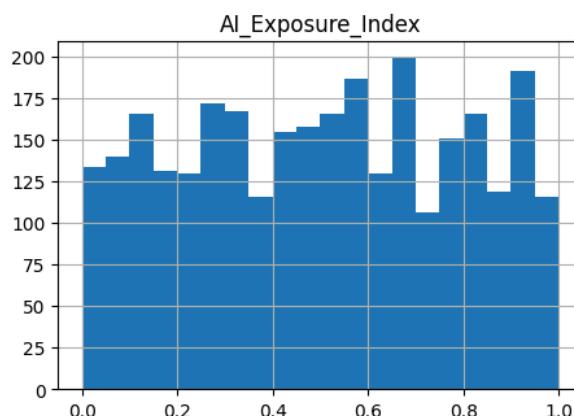
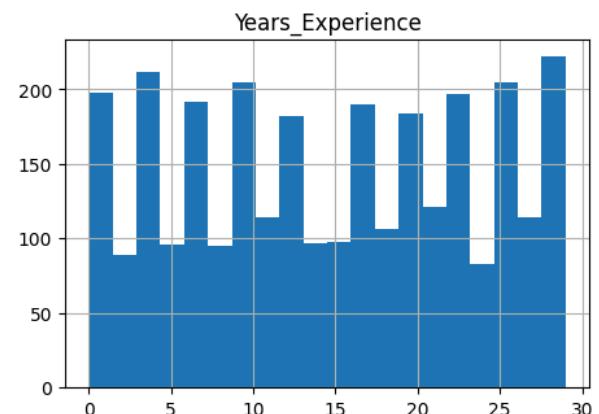
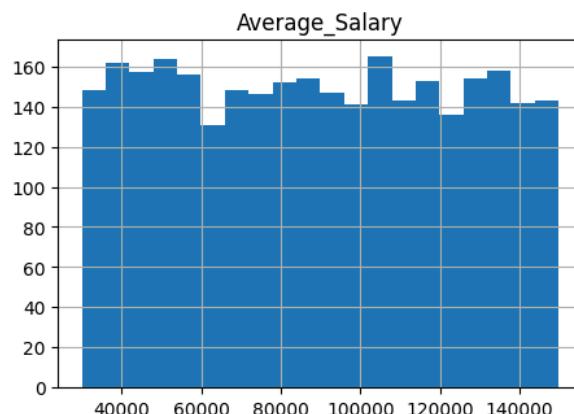
```
plt.figure(figsize=(12,6))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap="coolwarm")
plt.title("Correlation Heatmap")
plt.show()
```



```
numeric_cols = df.select_dtypes(include=np.number).columns[:4]

df[numeric_cols].hist(figsize=(12,8), bins=20)
plt.suptitle("Numeric Feature Distributions")
plt.show()
```

Numeric Feature Distributions



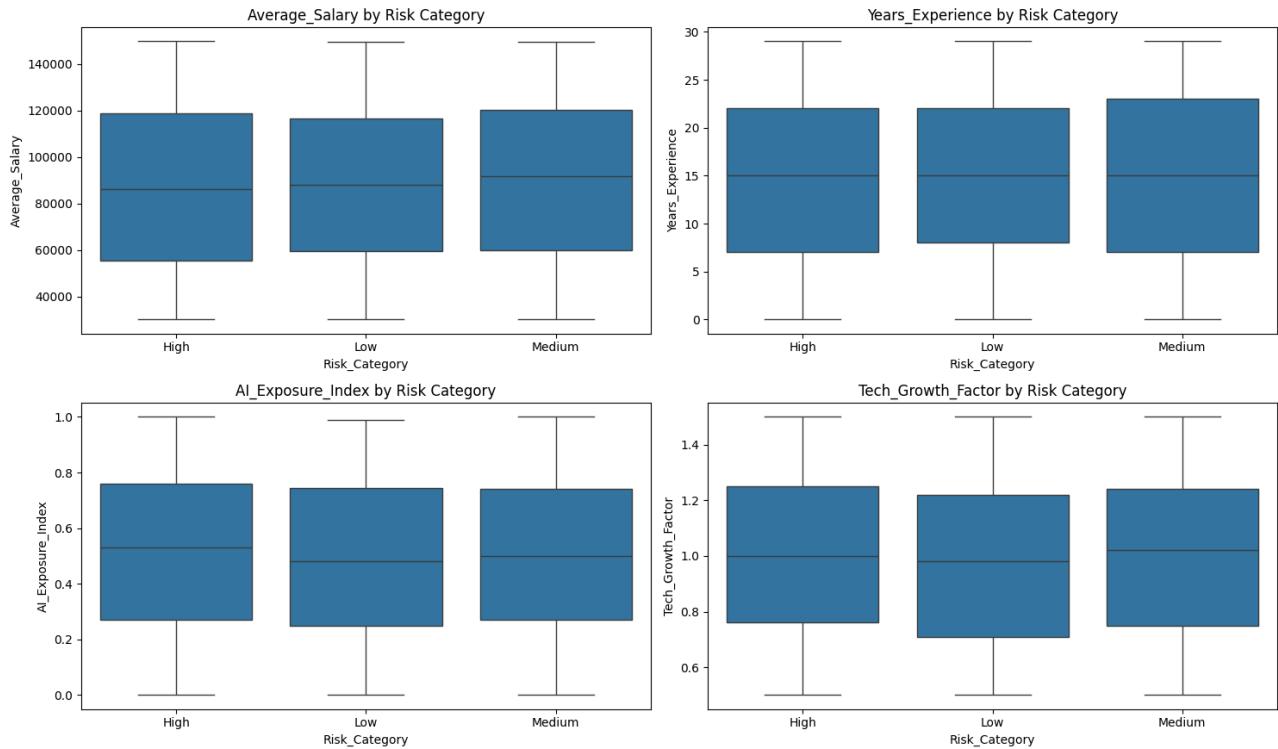
```
plt.figure(figsize=(15, 10))

numeric_cols_for_boxplot = ['Average_Salary', 'Years_Experience', 'AI_Exposure_Index', 'Tech_Growth_Factor']

for i, col in enumerate(numeric_cols_for_boxplot):
    plt.subplot(2, 2, i + 1) # Create a 2x2 grid for 4 plots
    sns.boxplot(data=df, x="Risk_Cat", y=col)
    plt.title(f'{col} by Risk Category')

plt.suptitle("Boxplots of Major Features by Risk Category", y=1.02) # Adjust suptitle position
plt.tight_layout(rect=[0, 0.03, 1, 0.95]) # Adjust layout to prevent title overlap
plt.show()
```

Boxplots of Major Features by Risk Category



```
df_encoded = pd.get_dummies(df, drop_first=True)
df_encoded.head()
```

	Average_Salary	Years_Experience	AI_Exposure_Index	Tech_Growth_Factor	Automation_Probability_2030	Sk
0	45795	28	0.18	1.28	0.85	
1	133355	20	0.62	1.11	0.05	
2	146216	2	0.86	1.18	0.81	
3	136530	13	0.39	0.68	0.60	
4	70397	22	0.52	1.46	0.64	

5 rows × 39 columns

```
from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()
df["Risk_Category_Encoded"] = label_encoder.fit_transform(df["Risk_Category"])
df["Risk_Category_Encoded"].head()
```

Risk_Category_Encoded

0	0
1	1
2	0
3	2
4	2

dtype: int64

```
y = df["Risk_Category_Encoded"]

X = df.drop(columns=["Risk_Category", "Risk_Category_Encoded"])
X = pd.get_dummies(X, drop_first=True) # encode remaining categorical features

X.head()
```

Average_Salary Years_Experience AI_Exposure_Index Tech_Growth_Factor Automation_Probability_2030 Sk

0	45795	28	0.18	1.28	0.85	
1	133355	20	0.62	1.11	0.05	
2	146216	2	0.86	1.18	0.81	
3	136530	13	0.39	0.68	0.60	
4	70397	22	0.52	1.46	0.64	

5 rows × 37 columns

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, random_state=42, stratify=y)
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
log_model = LogisticRegression(max_iter=1000, multi_class="auto")
log_model.fit(X_train_scaled, y_train)
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic.py:1247: FutureWarning: 'multi_class' warnings.warn(
```

```
    LogisticRegression
LogisticRegression(max_iter=1000, multi_class='auto')
```

```
y_pred = log_model.predict(X_test_scaled)
```

```
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

